

# Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\madrid_2003.csv")
df
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
0	01-03-2003 01:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.209999	NaN
1	01-03-2003 01:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.389999	NaN
2	01-03-2003 01:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.240002	NaN
3	01-03-2003 01:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.839996	NaN
4	01-03-2003 01:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.779999	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...
243979	01-10-2003 00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.380000	1.20
243980	01-10-2003 00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.400000	0.50
243981	01-10-2003 00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.830000	NaN

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY
243982	01-10-2003 00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.570000	NaN
243983	01-10-2003 00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.350000	2.43

243984 rows × 16 columns

## Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`In [4]: `df.columns`Out[4]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      33010 non-null   object 
 1   BEN        33010 non-null   float64
 2   CO         33010 non-null   float64
 3   EBE        33010 non-null   float64
 4   MXY        33010 non-null   float64
 5   NMHC       33010 non-null   float64
 6   NO_2       33010 non-null   float64
 7   NOx        33010 non-null   float64
 8   OXY        33010 non-null   float64
 9   O_3         33010 non-null   float64
 10  PM10       33010 non-null   float64
 11  PXY        33010 non-null   float64
 12  SO_2       33010 non-null   float64
 13  TCH        33010 non-null   float64
 14  TOL        33010 non-null   float64
 15  station    33010 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.3+ MB
```

In [6]: `data=df[['CO', 'station']]`  
`data`

Out[6]:

	CO	station
5	1.94	28079006
23	1.27	28079024
27	1.79	28079099
33	1.47	28079006
51	1.29	28079024
...	...	...
243955	0.41	28079099
243957	0.60	28079035
243961	0.82	28079006
243979	0.16	28079024
243983	0.29	28079099

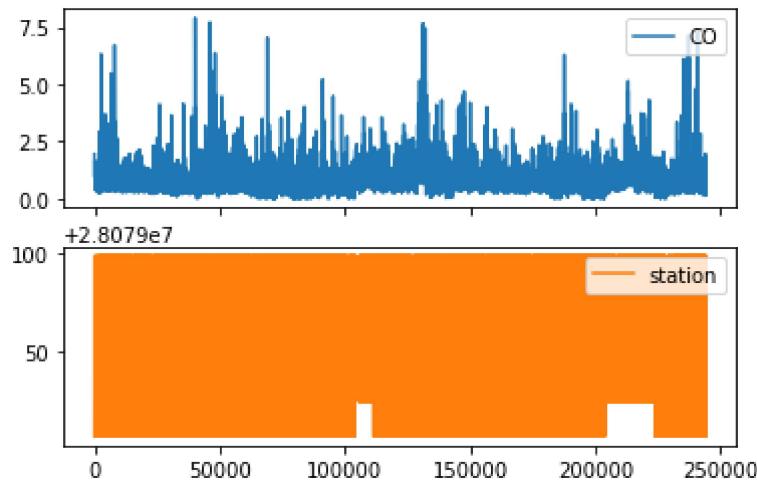
33010 rows × 2 columns

## Line chart

In [7]:

```
data.plot.line(subplots=True)
```

Out[7]: array([&lt;AxesSubplot:&gt;, &lt;AxesSubplot:&gt;], dtype=object)

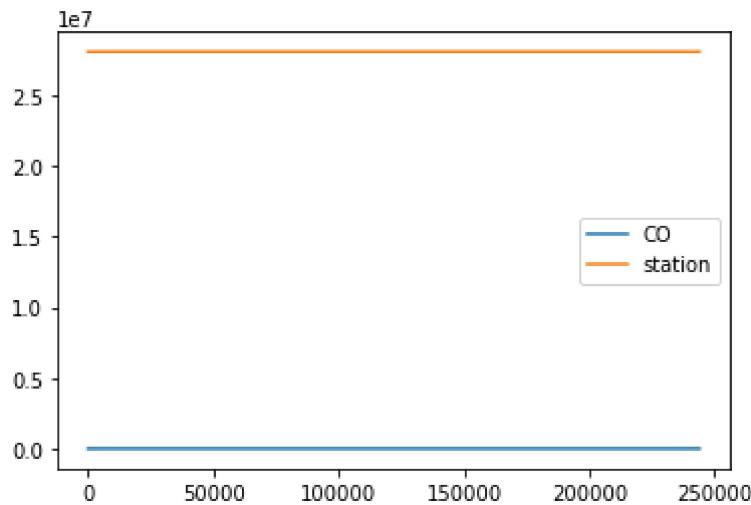


## Line chart

In [8]:

```
data.plot.line()
```

Out[8]: &lt;AxesSubplot:&gt;

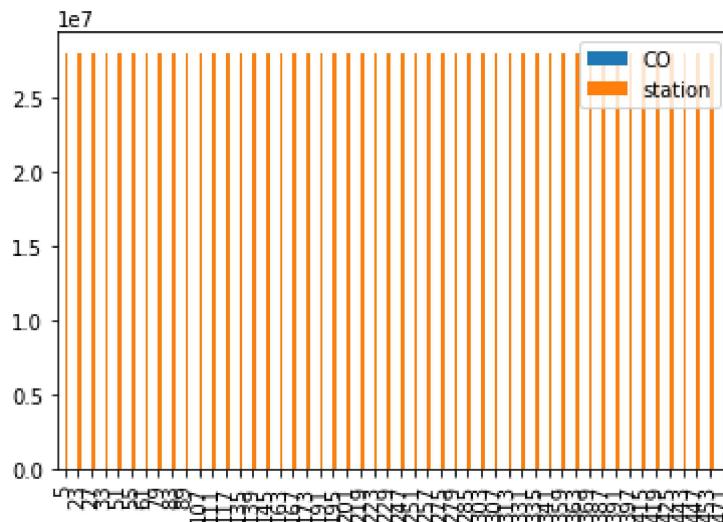


## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

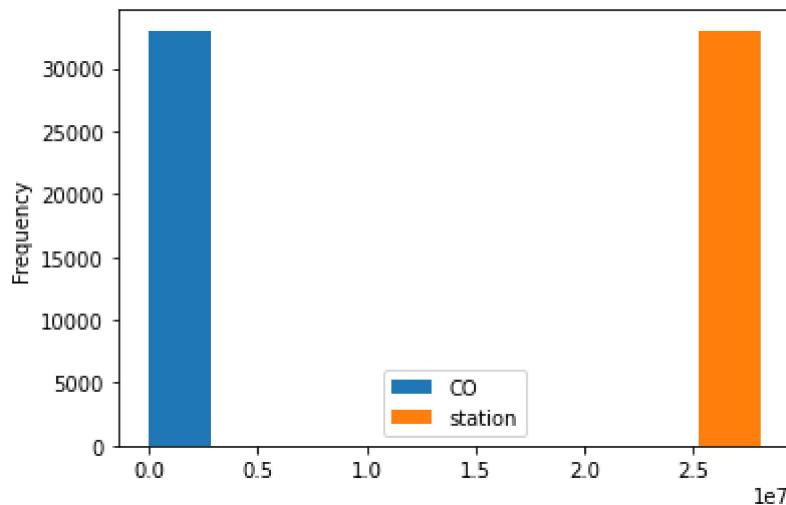
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```

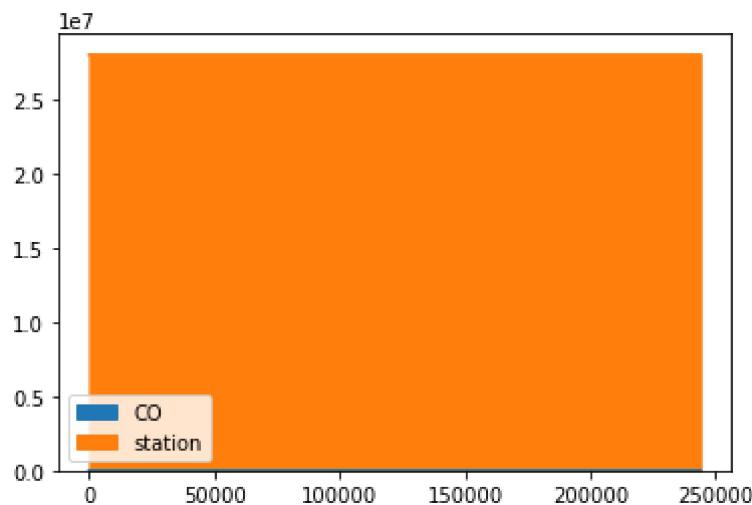
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: `data.plot.area()`

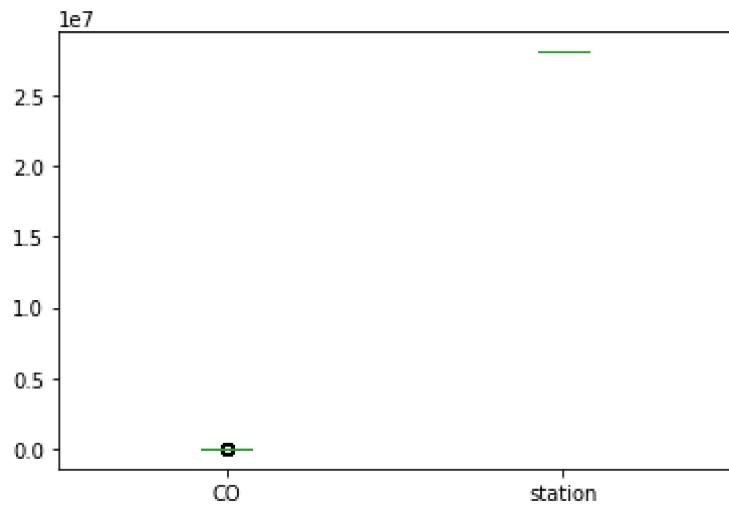
Out[12]: <AxesSubplot:>



## Box chart

In [13]: `data.plot.box()`

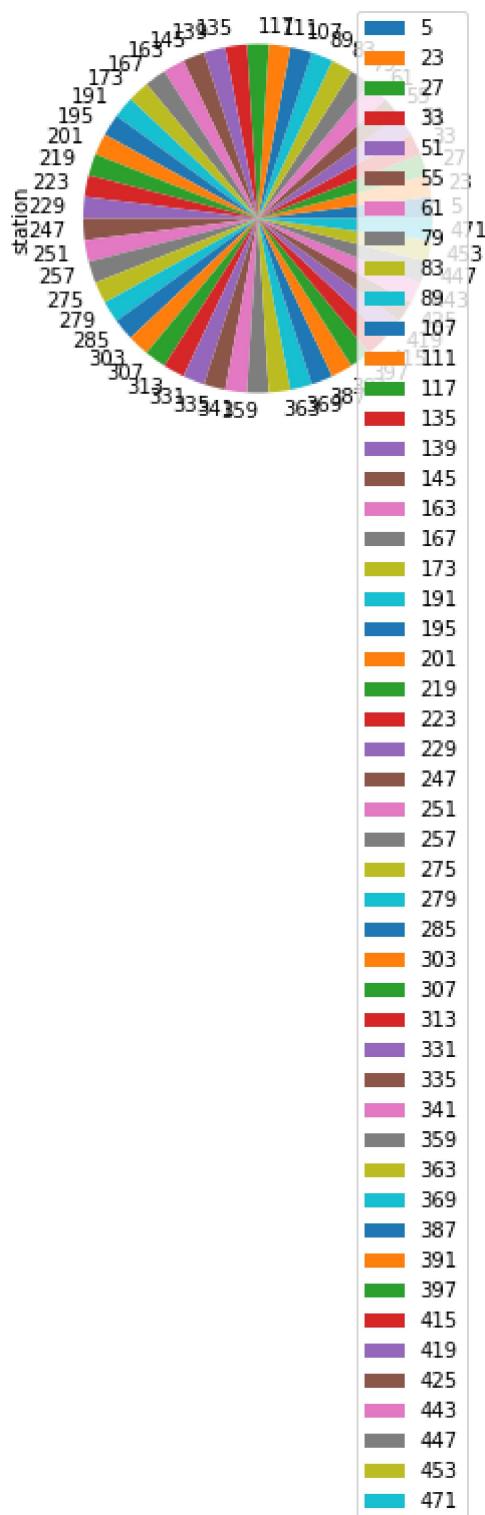
Out[13]: <AxesSubplot:>



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```

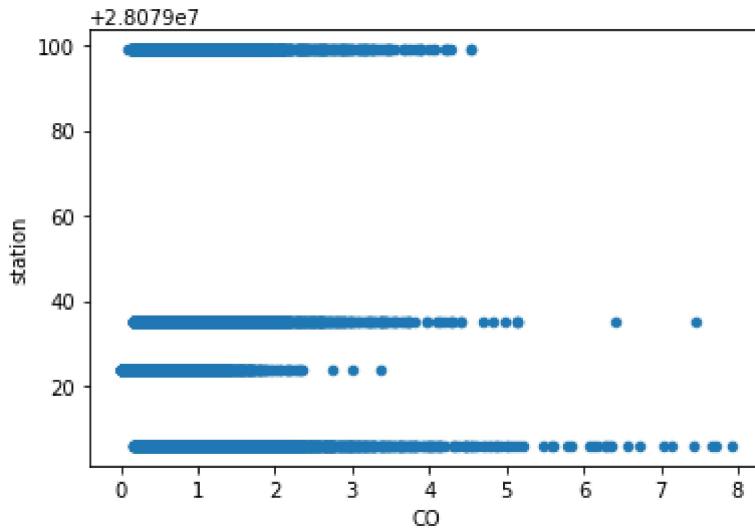
```
Out[14]: <AxesSubplot:ylabel='station'>
```



## Scatter chart

In [15]: `data.plot.scatter(x='CO' ,y='station')`

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        33010 non-null   object 
 1   BEN          33010 non-null   float64
 2   CO           33010 non-null   float64
 3   EBE          33010 non-null   float64
 4   MXY          33010 non-null   float64
 5   NMHC         33010 non-null   float64
 6   NO_2          33010 non-null   float64
 7   NOx          33010 non-null   float64
 8   OXY          33010 non-null   float64
 9   O_3           33010 non-null   float64
 10  PM10         33010 non-null   float64
 11  PXY          33010 non-null   float64
 12  SO_2          33010 non-null   float64
 13  TCH          33010 non-null   float64
 14  TOL          33010 non-null   float64
 15  station       33010 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.3+ MB
```

In [17]:

```
df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
<b>count</b>	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000	33010.000000
<b>mean</b>	2.192633	0.759868	2.639726	5.838414	0.137177	57.328049	120.153671
<b>std</b>	2.064160	0.545999	2.825194	6.267296	0.127863	31.811082	104.521701
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	0.900000	0.430000	1.010000	1.880000	0.060000	34.529999	49.070000
<b>50%</b>	1.610000	0.620000	1.890000	4.070000	0.110000	55.105000	92.779991
<b>75%</b>	2.810000	0.930000	3.300000	7.530000	0.170000	76.160004	160.100000

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
max	66.389999	7.920000	92.589996	177.600006	2.180000	342.700012	1246.000001

In [18]:

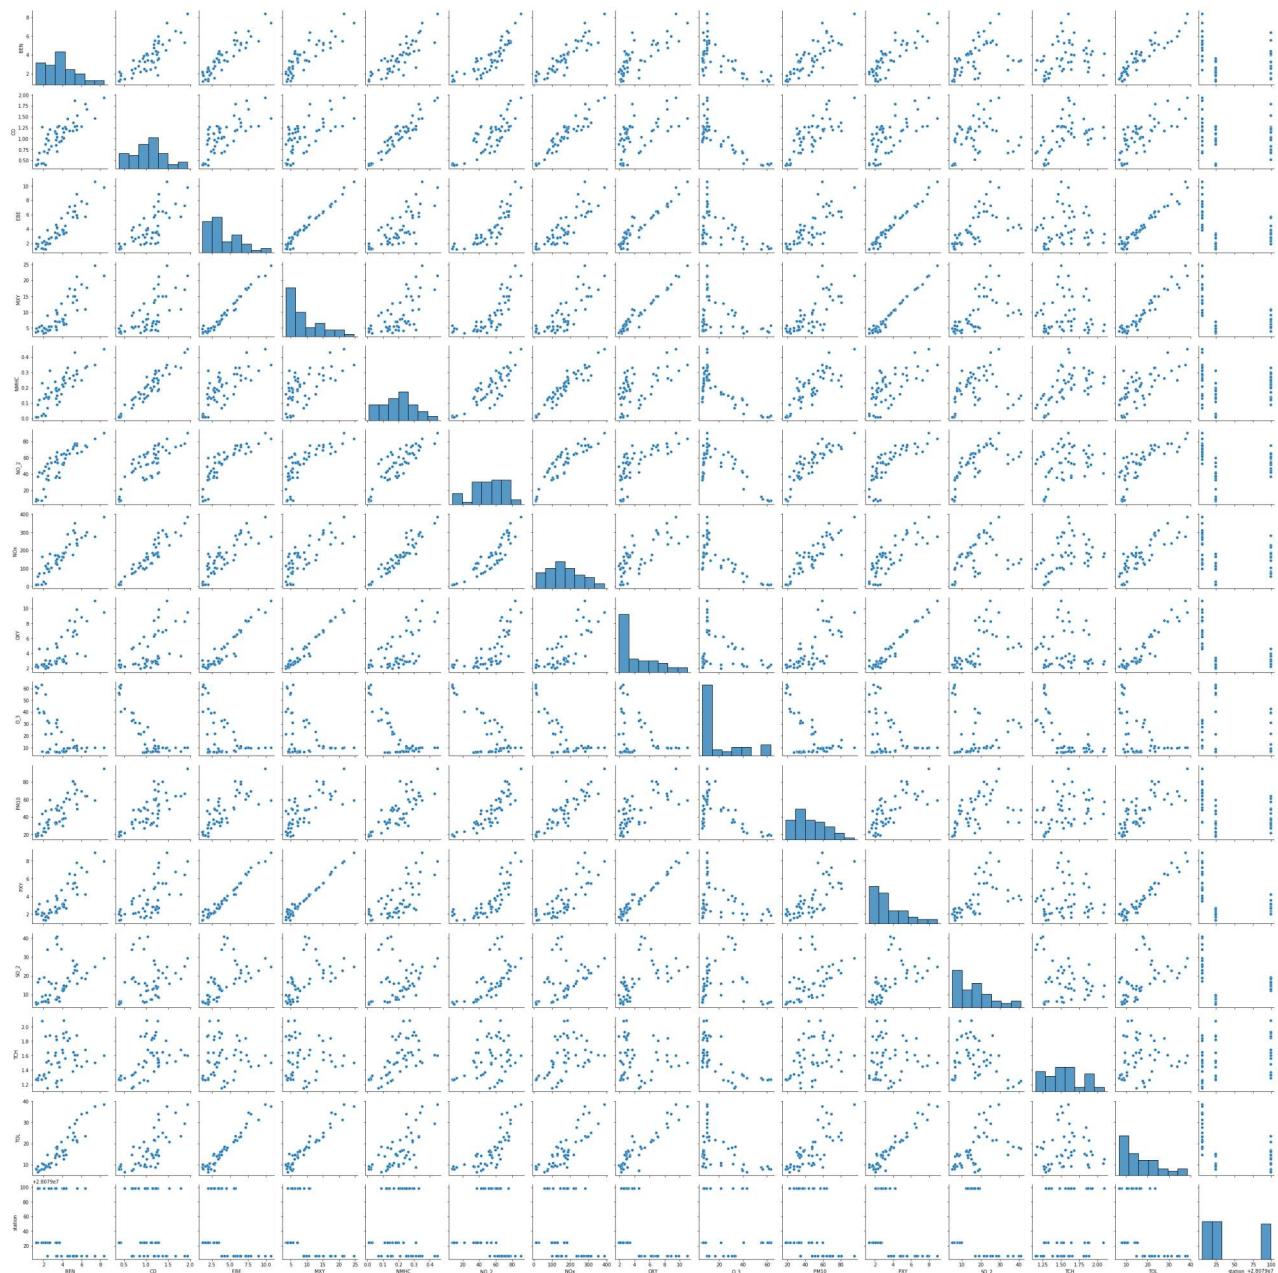
```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

## EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

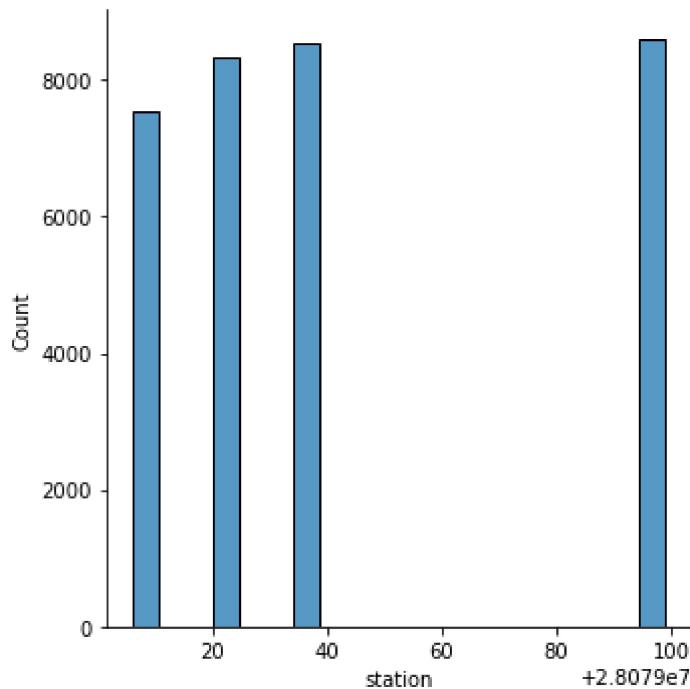
Out[19]:



In [20]:

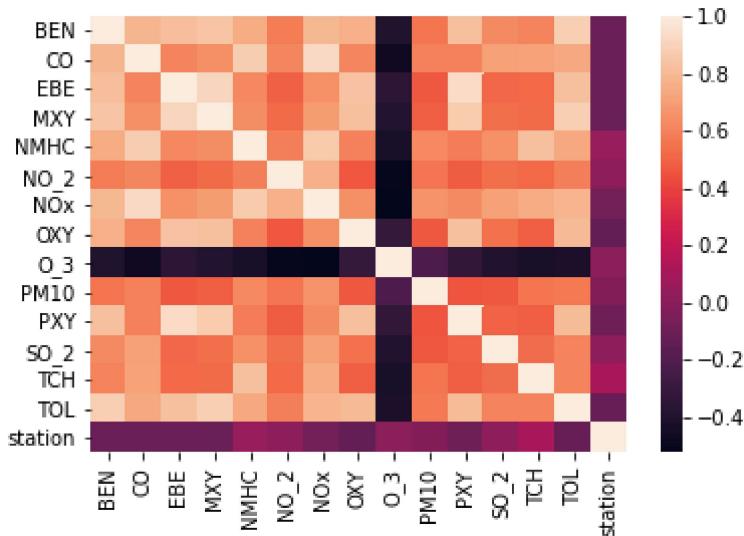
```
sns.displot(df1['station'])
```

Out[20]: <seaborn.axisgrid.FacetGrid at 0x198702d3ca0>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



## TO TRAIN THE MODEL AND MODEL BUILDING

In [22]:

```
x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

In [23]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28079001.705791198
```

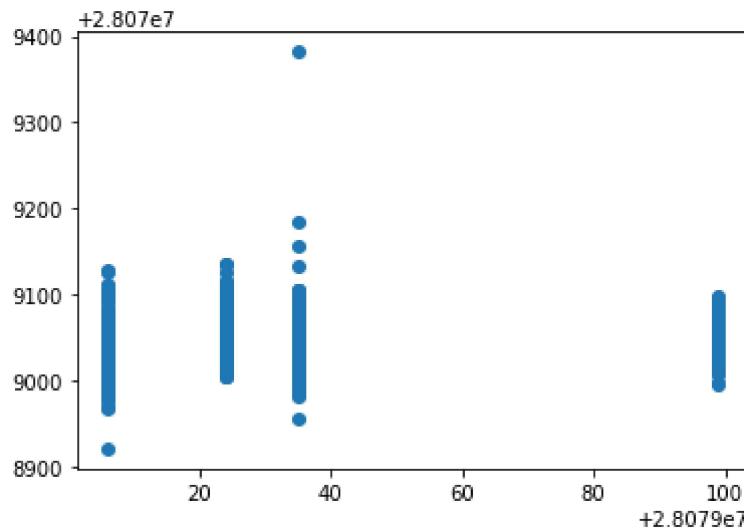
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
<b>BEN</b>	1.716726
<b>CO</b>	-40.340134
<b>EBE</b>	-1.691670
<b>MXY</b>	0.049299
<b>NMHC</b>	160.338149
<b>NO_2</b>	0.158254
<b>NOx</b>	-0.073205
<b>OXY</b>	-1.381130
<b>O_3</b>	-0.013500
<b>PM10</b>	-0.058948
<b>PXY</b>	1.865123
<b>SO_2</b>	0.847128
<b>TCH</b>	35.635134
<b>TOL</b>	-0.797461

```
In [27]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1987fa36730>
```



## ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.16672091906394193
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.1798509962774265
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.1674968474856896
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.1786612443473352
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[34]: Lasso(alpha=10)

```
In [35]: la.score(x_train,y_train)
```

Out[35]: 0.03629374637949523

## Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

Out[36]: 0.03451252116913206

## Elastic Net regression

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: en.coef_
```

Out[38]: array([ 0. , -0.30448294, 0.09232789, -0.11725017, 0.11276744,
 0.1526912 , -0.07341294, -1.20739406, -0.04598254, 0.08088606,
 0.30356586, 0.75909519, 1.54993948, -0.40026107])

```
In [39]: en.intercept_
```

Out[39]: 28079037.861311376

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

Out[41]: 0.047227086151283015

## Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
```

```
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

29.08043222878616  
1170.5345422485757  
34.213075603467395

## Logistic Regression

In [43]: `from sklearn.linear_model import LogisticRegression`

In [44]: `feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df['station']`

In [45]: `feature_matrix.shape`

Out[45]: (33010, 14)

In [46]: `target_vector.shape`

Out[46]: (33010,)

In [47]: `from sklearn.preprocessing import StandardScaler`

In [48]: `fs=StandardScaler().fit_transform(feature_matrix)`

In [49]: `logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)`

Out[49]: `LogisticRegression(max_iter=10000)`

In [50]: `observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]`

In [51]: `prediction=logr.predict(observation)  
print(prediction)`

[28079035]

In [52]: `logr.classes_`

Out[52]: `array([28079006, 28079024, 28079035, 28079099], dtype=int64)`

In [53]: `logr.score(fs,target_vector)`

```
Out[53]: 0.758467131172372
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.3461772071787448e-23
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[2.34617721e-23, 1.48032414e-55, 1.00000000e+00, 6.73763516e-16]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                 'min_samples_leaf':[5,10,15,20,25],
                 'n_estimators':[10,20,30,40,50]
                }
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters, cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                  'min_samples_leaf': [5, 10, 15, 20, 25],
                                  'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.7230709172578377
```

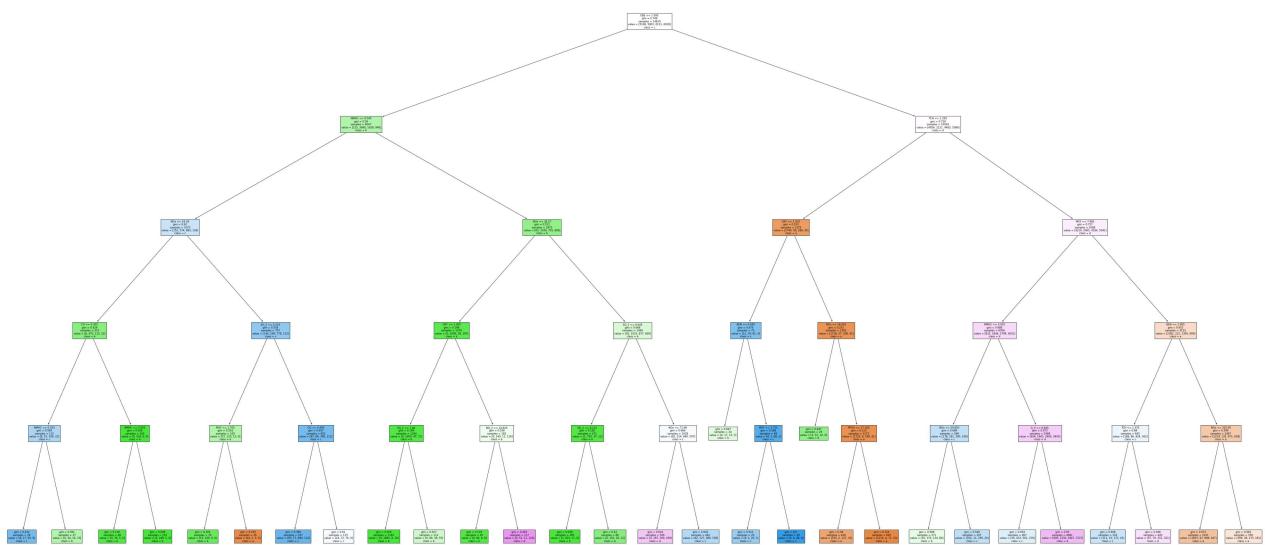
```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

```
Out[62]: [Text(2291.785714285714, 1993.2, 'EBE <= 1.095\ngini = 0.749\nsamples = 14610\nvalue = [5169, 5801, 6111, 6026]\nnclass = c'),
Text(1275.4285714285713, 1630.8000000000002, 'NMHC <= 0.045\ngini = 0.59\nsamples = 4047\nvalue = [215, 3680, 1628, 940]\nnclass = b'),
Text(637.7142857142857, 1268.4, 'NOx <= 19.26\ngini = 0.62\nsamples = 1072\nvalue = [152, 574, 893, 134]\nnclass = c'),
Text(318.85714285714283, 906.0, 'CO <= 0.305\ngini = 0.429\nsamples = 315\nvalue = [8, 375, 115, 22]\nnclass = b'),
Text(159.42857142857142, 543.5999999999999, 'NMHC <= 0.025\ngini = 0.584\nsamples = 115\nvalue = [8, 51, 109, 22]\nnclass = c'),
Text(79.71428571428571, 181.1999999999982, 'gini = 0.432\nsamples = 78\nvalue = [8, 17, 95, 9]\nnclass = c'),
Text(239.1428571428571, 181.1999999999982, 'gini = 0.591\nsamples = 37\nvalue = [0, 34, 14, 13]\nnclass = b'),
Text(478.2857142857142, 543.5999999999999, 'NMHC <= 0.005\ngini = 0.036\nsamples = 200\nvalue = [0, 324, 6, 0]\nnclass = b'),
Text(398.57142857142856, 181.1999999999982, 'gini = 0.116\nsamples = 48\nvalue = [0, 76, 5, 0]\nnclass = b'),
Text(558.0, 181.1999999999982, 'gini = 0.008\nsamples = 152\nvalue = [0, 248, 1, 0]\nnclass = b'),
Text(956.5714285714284, 906.0, 'SO_2 <= 5.215\ngini = 0.554\nsamples = 757\nvalue = [144, 199, 778, 112]\nnclass = c'),
Text(797.1428571428571, 543.5999999999999, 'MXY <= 1.705\ngini = 0.516\nsamples = 105\nvalue = [57, 110, 12, 0]\nnclass = b'),
Text(717.4285714285713, 181.1999999999982, 'gini = 0.293\nsamples = 79\nvalue = [13, 109, 9, 0]\nnclass = b'),
Text(876.8571428571428, 181.1999999999982, 'gini = 0.155\nsamples = 26\nvalue = [44, 1, 3, 0]\nnclass = a'),
Text(1116.0, 543.5999999999999, 'CO <= 0.495\ngini = 0.447\nsamples = 652\nvalue = [87, 89, 766, 112]\nnclass = c'),
Text(1036.2857142857142, 181.1999999999982, 'gini = 0.366\nsamples = 537\nvalue = [63, 17, 690, 112]\nnclass = c'),
Text(1195.7142857142856, 181.1999999999982, 'gini = 0.61\nsamples = 115\nvalue = [24, 72, 76, 0]\nnclass = c'),
Text(1913.1428571428569, 1268.4, 'NOx <= 36.27\ngini = 0.511\nsamples = 2975\nvalue = [63, 3106, 735, 806]\nnclass = b'),
Text(1594.2857142857142, 906.0, 'OXY <= 1.005\ngini = 0.198\nsamples = 1476\nvalue = [0, 2090, 58, 197]\nnclass = b'),
Text(1434.8571428571427, 543.5999999999999, 'SO_2 <= 7.48\ngini = 0.109\nsamples = 1294\nvalue = [0, 1950, 47, 71]\nnclass = b'),
Text(1355.142857142857, 181.1999999999982, 'gini = 0.026\nsamples = 1182\nvalue = [0, 1862, 9, 16]\nnclass = b'),
Text(1514.5714285714284, 181.1999999999982, 'gini = 0.627\nsamples = 112\nvalue = [0, 88, 38, 55]\nnclass = b'),
Text(1753.7142857142856, 543.5999999999999, 'NO_2 <= 14.825\ngini = 0.536\nsamples = 182\nvalue = [0, 140, 11, 126]\nnclass = b'),
Text(1673.999999999998, 181.1999999999982, 'gini = 0.118\nsamples = 65\nvalue = [0, 89, 0, 6]\nnclass = b'),
Text(1833.4285714285713, 181.1999999999982, 'gini = 0.483\nsamples = 117\nvalue = [0, 51, 11, 120]\nnclass = d'),
Text(2232.0, 906.0, 'SO_2 <= 6.645\ngini = 0.666\nsamples = 1499\nvalue = [63, 1016, 677, 609]\nnclass = b'),
Text(2072.5714285714284, 543.5999999999999, 'SO_2 <= 6.115\ngini = 0.126\nsamples = 471\nvalue = [1, 702, 37, 12]\nnclass = b'),
Text(1992.8571428571427, 181.1999999999982, 'gini = 0.038\nsamples = 385\nvalue = [1, 601, 11, 0]\nnclass = b'),
Text(2152.285714285714, 181.1999999999982, 'gini = 0.43\nsamples = 86\nvalue = [0, 101, 26, 12]\nnclass = b'),
Text(2391.428571428571, 543.5999999999999, 'NOx <= 71.96\ngini = 0.666\nsamples = 1028\nvalue = [62, 314, 640, 597]\nnclass = c'),
Text(2311.7142857142853, 181.1999999999982, 'gini = 0.624\nsamples = 566\nvalue = [0, 187, 260, 439]\nnclass = d'),
Text(2471.142857142857, 181.1999999999982, 'gini = 0.642\nsamples = 462\nvalue = [62, 127, 380, 158]\nnclass = c'),
Text(3308.142857142857, 1630.8000000000002, 'TCH <= 1.265\ngini = 0.729\nsamples = 1056
```

```
3\nvalue = [4954, 2121, 4483, 5086]\nclass = d'),
Text(2790.0, 1268.4, 'OXY <= 1.025\ngini = 0.313\nsamples = 1375\nvalue = [1740, 56, 28
9, 45]\nclass = a'),
Text(2630.5714285714284, 906.0, 'BEN <= 0.685\ngini = 0.474\nsamples = 73\nvalue = [12,
19, 81, 4]\nclass = c'),
Text(2550.8571428571427, 543.5999999999999, 'gini = 0.647\nsamples = 24\nvalue = [4, 1
7, 13, 3]\nclass = b'),
Text(2710.285714285714, 543.5999999999999, 'MXY <= 3.735\ngini = 0.248\nsamples = 49\nv
alue = [8, 2, 68, 1]\nclass = c'),
Text(2630.5714285714284, 181.1999999999982, 'gini = 0.512\nsamples = 20\nvalue = [8,
2, 20, 1]\nclass = c'),
Text(2790.0, 181.1999999999982, 'gini = 0.0\nsamples = 29\nvalue = [0, 0, 48, 0]\nclas
s = c'),
Text(2949.428571428571, 906.0, 'NOx <= 18.635\ngini = 0.252\nsamples = 1302\nvalue = [1
728, 37, 208, 41]\nclass = a'),
Text(2869.7142857142853, 543.5999999999999, 'gini = 0.447\nsamples = 29\nvalue = [3, 3
1, 10, 0]\nclass = b'),
Text(3029.142857142857, 543.5999999999999, 'PM10 <= 17.165\ngini = 0.223\nsamples = 127
3\nvalue = [1725, 6, 198, 41]\nclass = a'),
Text(2949.428571428571, 181.1999999999982, 'gini = 0.38\nsamples = 428\nvalue = [515,
2, 127, 31]\nclass = a'),
Text(3108.8571428571427, 181.1999999999982, 'gini = 0.124\nsamples = 845\nvalue = [121
0, 4, 71, 10]\nclass = a'),
Text(3826.2857142857138, 1268.4, 'MXY <= 7.885\ngini = 0.727\nsamples = 9188\nvalue =
[3214, 2065, 4194, 5041]\nclass = d'),
Text(3507.428571428571, 906.0, 'NMHC <= 0.055\ngini = 0.688\nsamples = 6056\nvalue = [8
32, 1844, 2799, 4051]\nclass = d'),
Text(3347.9999999999995, 543.5999999999999, 'NOx <= 59.855\ngini = 0.689\nsamples = 568
\nvalue = [178, 181, 390, 108]\nclass = c'),
Text(3268.285714285714, 181.1999999999982, 'gini = 0.668\nsamples = 271\nvalue = [16,
170, 134, 84]\nclass = b'),
Text(3427.7142857142853, 181.1999999999982, 'gini = 0.549\nsamples = 297\nvalue = [16
2, 11, 256, 24]\nclass = c'),
Text(3666.8571428571427, 543.5999999999999, 'O_3 <= 6.845\ngini = 0.673\nsamples = 5488
\nvalue = [654, 1663, 2409, 3943]\nclass = d'),
Text(3587.142857142857, 181.1999999999982, 'gini = 0.654\nsamples = 807\nvalue = [54,
433, 562, 216]\nclass = c'),
Text(3746.5714285714284, 181.1999999999982, 'gini = 0.65\nsamples = 4681\nvalue = [60
0, 1230, 1847, 3727]\nclass = d'),
Text(4145.142857142857, 906.0, 'BEN <= 2.985\ngini = 0.652\nsamples = 3132\nvalue = [23
82, 221, 1395, 990]\nclass = a'),
Text(3985.7142857142853, 543.5999999999999, 'TCH <= 1.375\ngini = 0.68\nsamples = 645\nn
value = [168, 86, 424, 362]\nclass = c'),
Text(3905.9999999999995, 181.1999999999982, 'gini = 0.608\nsamples = 202\nvalue = [11
1, 10, 172, 41]\nclass = c'),
Text(4065.428571428571, 181.1999999999982, 'gini = 0.648\nsamples = 443\nvalue = [57,
76, 252, 321]\nclass = d'),
Text(4304.571428571428, 543.5999999999999, 'NOx <= 320.95\ngini = 0.599\nsamples = 2487
\nvalue = [2214, 135, 971, 628]\nclass = a'),
Text(4224.857142857142, 181.1999999999982, 'gini = 0.573\nsamples = 1931\nvalue = [181
5, 97, 698, 447]\nclass = a'),
Text(4384.285714285714, 181.1999999999982, 'gini = 0.662\nsamples = 556\nvalue = [399,
38, 273, 181]\nclass = a')]
```



# Conclusion

## Accuracy

Linear Regression:0.1798509962774265

Ridge Regression:0.1786612443473352

Lasso Regression:0.1786612443473352

ElasticNet Regression:0.047227086151283015

Logistic Regression:0.758467131172372

Random Forest:0.7230709172578377

**Accuracy for logistic regression is higher so it is the best fit model**